

Development Roadmap M1–M5

UM-NATOS-007

DOCUMENT	REVISION	STATUS	CLASSIFICATION
UM-NATOS-007	1.0 · 2026-08-14	current	Internal

1. Abstract

This report defines the remaining milestones to a working operating system, each with its technical content, principal risk, and exit criteria. Milestones are ordered so that each one's correctness can be established before the next depends on it.

2. Milestone summary

ID	Deliverable	Principal risk	State
M0	Kernel boots, self-checks pass	Link map errors	Complete
M1	Timer interrupt and tick counter	Vector installation; silent faults	Next
M2	Two native tasks, preemptive switching	Context save correctness	—
M3	Heap allocator and VM memory model	Fragmentation; arena sizing	—
M4	Bytecode interpreter executing a program	Instruction set design	—
M5	Two VM applications time-sliced in bounded arenas	Isolation enforcement	—

3. M1 — Timer interrupt and tick counter

Deliverable. A periodic timer interrupt increments a kernel tick counter; `kmain` observes it advancing. The spin-loop heartbeat from M0 is replaced by a timed one.

Technical content

- Install an exception/interrupt vector table at the address the CPU expects (`VECBASE`).
- Write the interrupt entry and exit sequence. Under `call0` this is an explicit register save; the hardware does not do it.
- Configure a hardware timer and route it to a CPU interrupt level.
- Enable interrupts and confirm the handler runs without corrupting the interrupted context.
- Identify and either feed or disable any active watchdog.

Principal risk. A defective vector or entry sequence produces a reset with no output. The reset reason on the next boot is the only signal, and it is coarse. This is the milestone where the absence of JTAG hurts most.

Prerequisite — resolve UM-NATOS-004 §5. M1 will push `.text` past the point where the kernel's IRAM segment overlaps the bootloader's working memory. This must be addressed before flashing, or M1 will fail during load with symptoms resembling a code defect.

Exit criteria 1. Tick counter advances at a measured, stable rate. 2. Interrupted code resumes correctly (verified by a checksum over registers across an interrupt). 3. System survives ≥ 60 seconds without reset.

4. M2 — Native task switching

Deliverable. Two kernel tasks alternate under a preemptive scheduler driven by the M1 timer, each maintaining private state across switches.

Technical content

- Task control block: stack pointer, entry point, state.
- Per-task stack allocation with a guard pattern for overflow detection.
- Context switch in assembly: save `a0`, `a12 – a15` and live caller-saved registers, swap stack pointers, restore, return. **No window spilling** — see UM-NATOS-003 §6.
- Round-robin scheduler invoked from the timer interrupt.
- A yield primitive for cooperative switching.

Principal risk. An incorrect register save corrupts state in a way that manifests later, in unrelated code. Classic symptom: a task runs correctly for several switches then faults in a function that did nothing wrong.

Mitigation. Have each task write a distinctive register pattern before yielding and verify it on resume. This converts silent corruption into an immediate, localised assertion failure.

Exit criteria 1. Two tasks alternate for $\geq 10,000$ switches without corruption. 2. Register-pattern check passes on every resume. 3. Stack guard patterns intact.

5. M3 — Heap and VM memory model

Deliverable. A kernel allocator, and the arena model VM applications will run inside.

Technical content

- Allocator over the DRAM region. A simple free-list is adequate; predictability matters more than throughput.
- Fixed-size **arenas** for VM applications — contiguous blocks with recorded base and length, the bounds the interpreter will enforce.
- Arena lifecycle: allocate on application start, release on exit.
- Instrumentation: high-water mark, fragmentation measure, allocation failure counter.

Principal risk. Arena sizing. Too small and applications cannot do useful work; too large and few can run concurrently. With ~ 176 KB of DRAM total, minus kernel, stacks, and eventually a display buffer (UM-NATOS-

004 §4 notes a full 16-bit framebuffer alone would be ~150 KB), the budget is genuinely tight and should be settled with measurements rather than guesses.

Exit criteria 1. Allocate/free cycles leave no leak across 10,000 iterations. 2. Arena bounds are queryable by the interpreter. 3. Out-of-memory returns a failure rather than corrupting.

6. M4 — Bytecode interpreter

Deliverable. The VM executes a trivial program — arithmetic, a loop, a syscall producing UART output.

Technical content

- Instruction set design. Register-based is faster; stack-based is simpler to target. **Decision pending.**
- Dispatch loop with a per-instruction quantum counter — the mechanism enabling preemption at safe boundaries (UM-NATOS-001 §4.2).
- **Bounds-checked memory access.** Every load and store validated against the arena. This is the isolation guarantee; it is not optional and must not be compiled out for performance.
- Syscall instruction trapping into kernel services.
- A producer: assembler or compiler for the bytecode.

Principal risk. Instruction set design is difficult to revise once a producer and programs exist. Keep it small — roughly 40 opcodes — and resist convenience additions.

Secondary risk. Interpretation overhead is unmeasured. If it proves unacceptable, the response is a threaded-dispatch or computed-goto interpreter, not abandoning bounds checking.

Exit criteria 1. Program computing a known result returns it correctly. 2. Out-of-bounds access is refused and reported, not performed. 3. Interpreter yields at quantum expiry. 4. Instructions-per-second measured and recorded.

7. M5 — Multiple applications

Deliverable. Two VM applications run concurrently, time-sliced, each confined to its own arena, with a misbehaving application unable to affect the other.

Technical content

- Application table, scheduling, lifecycle.
- Per-application arena assignment and enforcement.
- Fault handling: an application violating its bounds is terminated, not the system.
- A minimal shell to start and stop applications.

Principal risk. This is where the isolation claim is tested. An application deliberately written to escape its arena must fail to do so.

Exit criteria 1. Two applications interleave correctly. 2. An application attempting out-of-bounds access is terminated; the other continues unaffected. 3. Terminating an application releases its arena completely.

8. Cross-cutting items

8.1 Version control — outstanding

Not yet initialised. The sibling project The Word Device was lost and recovered only because an editor's local history happened to retain it. This should be resolved before M1.

8.2 JTAG probe — ordered, not in hand

From M1 onward, failures are increasingly silent. UART cannot report a fault in the interrupt vector itself. Sequencing M1 after the probe arrives is recommended; if not, budget substantially more time.

8.3 Deferred decisions

Decision	Needed by	Notes
SMP or dedicated second core	M2	Fixes the scheduler's shape
Flash execute-in-place	When kernel exceeds IRAM	Depends on L1 cache configuration
Bytecode instruction set	M4	Expensive to revise afterwards
Filesystem strategy	Post-M5	Writing FAT is a substantial subproject
Display and touch drivers	Post-M5	Framebuffer memory is the binding constraint

8.4 Copy the borrowed binaries into this repository

`build.ps1` reads the bootloader and partition table from an unrelated project's build directory. Deleting or rebuilding that project breaks this one's flash step. Copying both into `nat-os/boot/` removes the coupling — about 21 KB.

9. References

- UM-NATOS-001 — architecture and the isolation model driving M4/M5
- UM-NATOS-003 §6 — why M2 is tractable
- UM-NATOS-004 §5 — the risk blocking M1
- UM-NATOS-006 — what M0 established