

# Milestone 4 Verification Report

UM-NATOS-012

| DOCUMENT     | REVISION         | STATUS                                                                               | CLASSIFICATION |
|--------------|------------------|--------------------------------------------------------------------------------------|----------------|
| UM-NATOS-012 | 1.1 · 2026-08-15 | <b>**PASS**</b> — all exit criteria met on hardware; §10 added, seven syscalls since | Internal       |

## 1. Abstract

Milestone 4 delivers the bytecode interpreter: a register machine with 16 registers and fixed 4-byte instructions, a host-side assembler, and syscalls into kernel services.

This is the milestone where nat-os acquires isolation. Everything before it was trusted code sharing one address space with nothing between components but convention. From here, application code runs under a bounds check on every memory access, and the measurements in §6 are the evidence that a malformed or hostile program stops itself rather than the kernel.

## 2. The decision: register-based

Recorded as pending in UM-NATOS-007 §6 and settled here.

|                                        | Stack machine                         | Register machine (selected)        |
|----------------------------------------|---------------------------------------|------------------------------------|
| Instructions per unit of work          | High — operands pushed and popped     | Low — operands named directly      |
| Encoding                               | Compact, often 1 byte                 | 4 bytes fixed                      |
| Decode                                 | Trivial                               | One byte plus three fields         |
| Producer complexity                    | Lower — expression trees map directly | Higher — needs register allocation |
| Dispatch overhead per useful operation | Paid several times                    | Paid once                          |

Dispatch is the inner loop of everything this kernel will ever run. A stack machine spends a large share of its instructions moving operands rather than computing with them, and each of those pays the full dispatch cost. A register machine executes the same program in materially fewer instructions, which more than repays a wider encoding.

The cost is a more demanding producer. That cost lands on the host, where there is no memory constraint and no consequence for being slow — see §4.

`sum(1..10)` compiles to **70 executed instructions** including the loop, syscalls, and a call/return.

## 3. Instruction set

35 opcodes, within the ~40 the roadmap allows. Fixed 4-byte encoding:

```

byte 0  opcode
byte 1  a      destination register, or the sole operand
byte 2  b      source register, or immediate low byte
byte 3  c      source register / offset, or immediate high byte

```

| Group              | Opcodes                                                 |
|--------------------|---------------------------------------------------------|
| Control            | HALT NOP MOV LDI LDIH                                   |
| Arithmetic / logic | ADD SUB MUL DIV MOD AND OR XOR SHL SHR SAR NOT NEG ADDI |
| Comparison         | SEQ SNE SLT SLTU SLE SLEU                               |
| Control flow       | JMP BRZ BRNZ CALL RET                                   |
| Memory             | LDW LDB STW STB                                         |
| Services           | SYS                                                     |

Comparisons produce 0 or 1 into a register rather than setting flags, so branches need only test a register for zero. That removes a whole class of state — condition codes with their own save/restore obligations across a context switch — from a kernel that has already paid once for forgetting to save processor state (UM-NATOS-009 §6).

### 3.1 Choices that close failure modes

- **Register indices are validated, not masked.** An index of 16 or more faults. Masking would turn a producer bug into wrong answers instead of a diagnosed stop.
- **Shift counts are masked to 5 bits.** A shift of 32 or more is undefined in C; leaving it undefined would let a program's behaviour depend on the host compiler, which is precisely what a VM exists to prevent.
- **Word accesses must be 4-byte aligned.** Faulting is deterministic; the alternative is an unaligned-access exception in kernel context.
- **Branch displacements count instructions,** so a branch cannot target a misaligned address.

### 3.2 Return addresses live in kernel memory

`CALL` pushes to a kernel-side stack of 32 entries in the `vm_t` structure, not into the arena.

A program therefore **cannot address its own return addresses**. No amount of wild pointer arithmetic within its arena can corrupt where a function returns to, because that data is not in the arena at all. The entire family of stack-smashing bugs is unexpressible rather than defended against.

The cost is a fixed call depth and no stack-allocated locals in the conventional sense. For a kernel whose isolation is entirely software-enforced, being unable to express the attack is worth more than the flexibility.

## 4. The producer

`tools/vasm.py` assembles a text source to a C header containing a byte array and label offsets. It supports labels, `.string`, `.byte`, `.word`, `.align`, `.space`, character literals, `@label` for addresses, and instruction-relative branch resolution.

It runs on the **host**. Nothing about a producer needs to live in 176 KB of DRAM, and keeping it off the device means the on-device side stays a pure interpreter with no parsing, no symbol table, and no attack surface from untrusted text. `build.ps1` assembles `tools/*.vasm` into `kernel/generated/` as a build product.

The demonstration program is 120 bytes and 6 labels.

## 5. Preemption boundary

`vm_run(vm, quantum)` executes at most `quantum` instructions and returns `VM_RUN_QUANTUM` if the budget runs out. State lives entirely in `vm_t`, so execution resumes exactly where it stopped.

This is the safe-boundary preemption of UM-NATOS-001 §4.2: a task hosting a VM can be descheduled between any two bytecode instructions without the program being able to observe it. The interpreter never has to be reentrant, and the timer interrupt never has to understand bytecode.

A program that never terminates costs its quantum and nothing further — proved in §6.3 rather than asserted.

## 6. Results

```
vm says sum(1..10) = 55
[4a] program : HALTED ok  insns=70 resumes=4 status=0 fault=none
[4b] faults  : bounds=ok div0=ok opcode=ok reg=ok ret=ok align=ok PASS
[4c] runaway : PASS  bounded at 500 insns, no fault, kernel alive
[4d] predicate: PASS  vm_in_bounds agrees with arena_contains over 35 cases
```

### 6.1 The program

Arithmetic, a counted loop, a bounds-checked store and reload through arena memory, a call and return, and four syscalls producing UART output. It ran to `HALT` with exit status 0 in 70 instructions.

It was run with a quantum of 16, so it was **suspended and resumed 4 times** mid-execution and produced the correct answer regardless — which is the property that matters, not the answer itself.

### 6.2 Containment

Six deliberately malformed programs, each hand-encoded as raw bytes rather than assembled. Routing them through the assembler would only have demonstrated that well-formed programs behave; the claim under test is about malformed ones.

| Probe  | Attempts                                            | Expected fault               | Result |
|--------|-----------------------------------------------------|------------------------------|--------|
| bounds | Store to offset <code>0xFFF0</code> in a 2 KB arena | <code>VM_FAULT_BOUNDS</code> | ok     |
| div0   | Divide by zero                                      | <code>VM_FAULT_DIV0</code>   | ok     |
| opcode | Byte <code>0xFF</code> as an instruction            | <code>VM_FAULT_OPCODE</code> | ok     |
| reg    | <code>MOV r16, r0</code>                            | <code>VM_FAULT_REG</code>    | ok     |
| ret    | <code>RET</code> with an empty call stack           | <code>VM_FAULT_RET</code>    | ok     |
| align  | Word load at offset 1                               | <code>VM_FAULT_ALIGN</code>  | ok     |

Each stopped its own program with the correct code and returned control normally. The strongest evidence is indirect: every line of output after this test exists only because the kernel survived all six.

### 6.3 Runaway

A one-instruction program branching to itself. Executed exactly 500 instructions under a 500-instruction quantum, reported `VM_RUN_QUANTUM`, raised no fault, and left the kernel running. An unterminating program is a scheduling cost, not a hang.

### 6.4 The two bounds predicates agree

`vm_in_bounds()` duplicates `arena_contains()` so the interpreter's inner loop need not make a function call per access. Duplicated logic drifts, so the agreement is **tested rather than trusted**: 35 combinations of offset and length, including zero lengths, the exact end, one past the end, and lengths chosen to wrap the address space. All agree.

### 6.5 Regression

M3 self-test passing in full. M2 workload unchanged: 3,418 ticks, switches 1140/1139/1139, guards intact, `corrupt=0`.

### 6.6 Running under the scheduler

The VM now runs as a native task alongside the M2 workload rather than from the boot path, which is what makes the two preemption mechanisms observable together.

```
vm arena      : id=0 base=0x3ffb23f0 program=28 B
tasks         : report=0 a=1 b=2 vm=3

t=1801 switches r/a/b=451/450/450 work a/b=10585776/11610220 guards=ok
      freew a/b=480/480 corrupt=0
      | vm sw=450 insns=3291313 counter=1097103 fault=none
```

The hosted program increments a counter in its own arena forever, performing a bounds-checked store every iteration, so the isolation path is on the hot loop rather than exercised once at startup.

**The two mechanisms compose.** The timer interrupt suspends the VM task wherever it happens to be — almost always somewhere inside the interpreter's own C code, mid-instruction from the bytecode's point of

view — and `vm_run()` separately returns at a bytecode instruction boundary when its quantum expires. Neither is aware of the other, and the program cannot observe either.

**The arithmetic is the proof.** The loop body is exactly three instructions ( `add` , `stw` , `jmp` ), so instructions retired must be three times the counter plus the three-instruction preamble:

```
1,097,103 × 3 + 3 = 3,291,312    observed 3,291,313
```

The residual of one is the sample being taken mid-iteration, between the increment and the store. Across **450 preemptions** and 3.29 million bytecode instructions, not one instruction was lost, replayed, or half-executed. A context switch that dropped or repeated an interpreter step would show here as drift, and there is none.

**Scheduling is fair.** 451/450/450/450 across four tasks. The VM is an ordinary task with no special standing, and a program that never terminates costs exactly its share.

## 6.7 Dispatch cost

Derived from the same run, and worth recording because §9 previously listed it as unmeasured.

Between consecutive reports 200 ticks apart, the VM retired 365,704 instructions. 200 ticks is 160,000,000 CCOUNT cycles, of which the VM task receives roughly one quarter under an even four-way round robin:

```
40,000,000 cycles ÷ 365,704 instructions ≈ 109 cycles per bytecode instruction
```

That figure covers full dispatch, operand decode, register-range validation, and — for two of the three instructions in the loop — a bounds-checked memory access. It is a ceiling rather than a precise cost, since the quarter-share assumption ignores time spent in the interrupt handler itself.

109 cycles is high enough that a computed-goto dispatch table would be worth measuring if the VM ever becomes the bottleneck, and low enough that nothing here justifies trading away the diagnosability of a `switch` today.

## 7. `kstring.c` , and a trap avoided

The link failed on an undefined reference to `memcpy` from a function whose source never mentions it. GCC synthesises calls to `memcpy` / `memset` from ordinary C — byte-copy loops, struct assignments, large local initialisers — and does so even under `-fno-builtin` , which only governs the treatment of those names when written explicitly. Under `-nostdlib` nothing supplies them.

`kernel/kstring.c` provides `memcpy` , `memset` , `memmove` and `memcmp` .

The trap worth recording is what happens next: GCC will recognise the copy loop *inside* `memcpy` as a `memcpy` and rewrite it into a call to itself. That bug is silent, infinitely recursive, and would surface as a stack overflow in whatever unrelated code first copied a struct. The build now passes `-fno-tree-loop-distribute-patterns` , which is the supported way to prevent it.

## 8. Metrics

| Quantity                                        | Value                             |
|-------------------------------------------------|-----------------------------------|
| Opcodes                                         | 35 (roadmap ceiling ~40)          |
| Instruction width                               | 4 B fixed                         |
| Registers                                       | 16                                |
| Call depth                                      | 32, in kernel memory              |
| Demo program                                    | 120 B, 6 labels                   |
| Instructions to compute <code>sum(1..10)</code> | 70                                |
| Quantum resumptions during demo                 | 4                                 |
| Fault classes exercised                         | 6 of 10 defined                   |
| Bounds predicate cases cross-checked            | 35                                |
| Bytecode instructions under the scheduler       | 3,291,313                         |
| Preemptions survived                            | 450                               |
| Instruction-accounting drift                    | 1 (mid-iteration sample)          |
| Dispatch cost                                   | ~109 cycles/instruction (ceiling) |
| Image size                                      | 10,560 B                          |

## 9. What M4 does not establish

- **No arena-relative program loading.** A program is copied to offset 0 and addresses everything relative to its arena. There is no relocation, no linking of separately assembled units, and no code/data separation within the arena — a program can overwrite its own instructions.
- **Only 6 of 10 fault classes are exercised.** `VM_FAULT_PC`, `VM_FAULT_CALL_DEPTH`, `VM_FAULT_SYSCALL` and `VM_FAULT_STRING` are implemented but untested on hardware.
- **Dispatch cost is bounded, not precisely measured.** §6.7 derives ~109 cycles per instruction from a quarter-share assumption that ignores time spent in the interrupt handler, so it is a ceiling. A direct measurement would need `CCOUNT` sampled either side of a `vm_run()` call.
- **No multiple concurrent VMs.** One at a time, though nothing in the design prevents more — `vm_t` carries all state, and a second would need only a second arena and task.
- **The hosted program is not replaced when it stops.** A halted or faulted VM yields its task forever rather than loading something else; there is no program loader, only a copy into an arena at boot.
- **Syscall argument validation is per-call, not systematic.** Every syscall added since checks its own arguments (§10.1), and each was reasoned about individually. Nothing enforces that a future one will, and there is no shared harness that would catch an unchecked length in a new service.

## 10. Syscalls added after M4

Revision 1.1. M4 shipped five syscalls, all of them producing UART output or returning a scalar. Seven more have been added since, and they fall into three groups by what they hand across the boundary.

| Syscall                   | Added in        | Crosses the boundary                                         |
|---------------------------|-----------------|--------------------------------------------------------------|
| EXIT PUTC PUTS PUTD TICKS | M4              | Scalars, and one string read out of the arena                |
| FILL TEXT DIMS            | UM-NATOS-016    | Coordinates, clipped to a viewport                           |
| TOUCH                     | UM-NATOS-017 §8 | Coordinates, <i>inward</i> , filtered by viewport            |
| BLIT                      | UM-NATOS-016    | A pointer <b>and</b> a <b>length</b> , both program-supplied |
| SEND RECV                 | UM-NATOS-013 §8 | A buffer, copied through a kernel mailbox                    |

The ISA itself is unchanged: `SYS` is still one opcode of 35, and every addition is a service number rather than an instruction. That was the point of spending an opcode on a dispatch number instead of encoding services directly — the instruction set has not had to move since it was fixed.

### 10.1 Each group needed a different check

**Coordinates** are clipped in the offset domain, so a value near `0xFFFFFFFF` — which is what a program passing a negative number actually sends — fails the comparison rather than wrapping into range.

**BLIT** is the first syscall taking a pointer and a length together, so it is the first where the size is program-controlled. `w * h * 2` on unchecked 32-bit values overflows readily;  $65536 \times 65536$  wraps to zero, which would satisfy a byte-length check and then copy whatever followed the buffer. The dimensions are therefore each bounded against the panel *before* they are multiplied, after which the product provably cannot wrap.

**SEND** and **RECV** check the buffer in `vm.c`, not in `ipc.c`. Only the VM knows which arena an offset belongs to; the messaging layer receives a kernel pointer and has no way to tell where it came from. Putting the check where the information is, rather than where the operation happens, is the whole reason that split exists.

### 10.2 The quantum distinction

**FILL**, **TEXT** and **BLIT** end the caller's time slice; **TOUCH**, **SEND** and **RECV** do not.

The rule is the cost of the work rather than the fact of being a syscall. A fill is one instruction and milliseconds of SPI, so without ending the slice a 2,000-instruction quantum becomes minutes of drawing (§5). **TOUCH** reads a snapshot the polling task already published and costs nothing, so charging it a slice would only make an application slower for asking.

## 11. References

- UM-NATOS-001 §4.2 — isolation model and safe-boundary preemption
- UM-NATOS-007 §6 — M4 deliverable and the ISA decision recorded as pending
- UM-NATOS-010 §5.2 — `arena_contains()` and the offset-domain argument
- UM-NATOS-009 §6 — the saved-state failure that motivated §3 flag-free design
- `kernel/vm.h` — ISA definition and isolation rationale

- `kernel/vm.c` — dispatch loop and checked accessors
- `tools/vasm.py` — the producer